

The Competitive Landscape for AI-Powered Go Game Review & Commentary Tools (May 2026)

TL;DR

- **There is a real, unfilled product gap in natural-language Go commentary.** Every shipping Go review tool — AI Sensei, KaTrain, OGS, ZBaduk, Lizzie, GoReviewPartner — outputs only win-rate graphs, point-loss numbers, color-coded suggestions, and variation trees. None explain in plain language *why* a move is good or bad. Chess.com's "Game Review with Coach Explanations" (launched with the September 1, 2022 membership restructure; Diamond tier ~\$10/mo annualised, ~\$19.99/mo monthly) is the proven precedent; nothing equivalent exists for Go.
- **The market is mid-sized but underserved:** estimates of Go players worldwide range from ~24 million (Japanese census, 2002) to 60+ million (British Go Association, 2016), heavily concentrated in China (~40 M players per the Chinese Weiqi Association via CNN 2023), Korea, and Japan. The English-speaking, OGS-using subset is smaller but disproportionately willing to pay — Google DeepMind's official AlphaGo page states the Lee Sedol match was "*watched by over 200 million people worldwide*," and that 2016 awareness shock continues to feed the review-tool market. (Wikipedia + 3)
- **OGS's third-party extension ecosystem is essentially empty.** Greasy Fork lists only 7 userscripts for online-go.com — all cosmetic UI tweaks with single-digit install counts — and no commentary overlay or LLM-powered review extension has ever shipped. The OGS REST + Socket.IO API is documented and open, making this a low-friction technical wedge.

Key Findings

1. **The dominant paid SaaS in the category is AI Sensei** (ai-sensei.com), built on KataGo, founded by Benjamin Teuber (6-dan, German champion) with Erik Ouchterlony, Max Dohme, and Vitali Henne. Four tiers (Free 50 playouts → Basic 500 → Dan 2,500 → Pro 10,000 playouts/move). It still does not output any natural-language commentary — only point-loss/winrate visualizations and variation trees. (Ai-sensei) (Ai-sensei)
2. **The dominant free/open-source tool is KaTrain** (Sander Land, (sanderland/katrain) on GitHub), bundled with KataGo, distributed via PyPI / Homebrew / SourceForge. Focused on point-loss color-coded dots and humanlike weakened bots for teaching games. No NL commentary. (GitHub)
3. **OGS has a built-in AI review (Leela Zero & KataGo)** that runs free on all finished 19×19/13×13/9×9 games; site supporters get full per-move analysis. It is the highest-volume review surface in the Western Go world but offers only graphs and dot markers — community feedback explicitly asks for explanations. (Online Go Forum) (GitHub)
4. **GoReviewPartner** ((pnprog/goreviewpartner) on GitHub) **is essentially dormant** — last release v0.14.2 in July 2018, Python 2 codebase. It supports Leela Zero, AQ, Leela, Pachi, PhoenixGo, Ray, GnuGo. Useful primarily for offline batch analysis. (Yuntingdian) (GitHub)

5. **Sabaki** ([SabakiHQ/Sabaki](#)) is the **de facto open-source SGF editor**, free, Electron-based, supports GTP engines via plug-in. Not an AI reviewer itself; it's the "VS Code" of Go board editors. [Yichuanshen](#)
6. **Chess.com's "Game Review with Coach Explanations" is the validated precedent**. It generates text descriptions ("This defends a piece that was under attack...") gated behind Platinum (graphs/classifications) and Diamond (explanations) tiers. Per Chess.com's own help center: *"Platinum members have unlimited access to game reviews, while Diamond members also receive unlimited access to Coach explanations."* Diamond costs roughly \$10/mo annualised (\$120/yr) or ~\$19.99/month billed monthly. Lichess offers only the engine-based analysis without natural-language commentary as of May 2026. ChessCoach ([chrisbutner/ChessCoach](#)) is the only neural-net chess engine with NL commentary and runs as @PlayChessCoach on Lichess — proof that the technique is feasible. [chess + 2](#)
7. **Community frustration with "the AI doesn't explain why" is documented and consistent** across OGS forums, Life in 19x19, and AI Sensei's own blog. AI Sensei explicitly acknowledges in its own marketing: *"AI won't explain why a move is good in human terms. It won't say 'this move is important because it creates influence toward the center while limiting your opponent's base.'"* [AI-sensei](#)
8. **No OGS commentary extension has ever shipped**. All 7 Greasy Fork userscripts for online-go.com are cosmetic UI tweaks (move counter, board resizer, SGF copy button) with installs in single digits. This is open territory.

Details

Existing Competitors and Tools

AI Sensei ([ai-sensei.com](#)) — Closed-source SaaS, KataGo cloud backend. Free tier (50 playouts/move, ~4-move variation depth); Basic (500 playouts, ~10-move variations, KataGo's strength described as "superhuman"); Dan (2,500 playouts, ~15-move variations); Pro (10,000 playouts, ~18-move variations). Accepts SGF, GIB (Tygem), NGF (WBaduk), UGF (Pandamet), UGI. Permanent storage of reviews, shareable links, integrated with Pandamet/GoPanda2 ("Analyze with AI" post-game button), active Discord and YouTube channel. **Weaknesses:** no NL commentary, no in-game live review, free tier may misread life-and-death situations. Founder Benjamin Teuber stated to USGA: *"I believe the key feature of this tool is the focus on big mistakes; that way you can work on your biggest weaknesses instead of getting lost in minor details."* No public subscriber count is disclosed. [AI-sensei + 7](#)

KaTrain ([sanderland/katrain](#)) — Free, open-source (MIT-style), Python/Kivy, pre-bundled with KataGo (OpenCL build). Distributed for Windows, macOS (Homebrew), Linux (pipx). Marquee features: colored point-loss dots on each move; weakened-strength bots calibrated to ranks; teaching mode that auto-undoes bad moves; distributed-training contribution mode (KaTrain 1.8.0+). Best for non-technical users running local AI analysis on their own hardware. **Weaknesses:** requires a reasonable GPU for strong analysis; UI is functional but dated; no NL commentary; no cloud storage or sharing. [GitHub + 3](#)

Sabaki ([SabakiHQ/Sabaki](#)) — Free, open-source, Electron, cross-platform. SGF editor with markup tools, Markdown comments, game-tree navigation, score estimator, GTP engine integration with KataGo/Leela. Not an AI reviewer in its own right — it's the canonical SGF authoring environment used by reviewers, teachers, and content creators. Often paired with KataGo/Lizzie. [Yichuanshen](#) [Medium](#)

GoReviewPartner ([pnprog/goreviewpartner](#)) — Free, GPLv3, Python 2 (last release v0.14.2 on 2018-07-28). Supports Leela Zero, AQ, Leela, Pachi, PhoenixGo, Ray, GnuGo. Designed for offline batch analysis of SGFs. **Effectively legacy** today; users have migrated to KaTrain or AI Sensei.

[GitHub](#)

OGS built-in AI Review ([online-go.com](#)) — Free Leela Zero analysis for every finished game; site supporters get full KataGo per-move review and interactive variation analysis (announced as "Interactive AI analysis is now available" by anoek). Win-rate bar, score-difference bar, key-move list, "AI suggested" variation playback. The user-facing wiki ("Understanding the AI review") documents what each color/dot means but doesn't generate text per move. The OGS forum has multiple long threads from users asking for explanations beyond percentages. [Online Go Forum + 2](#)

Leela Zero and KataGo (engines) — KataGo ([lightvector/KataGo](#)) is the strongest open-source Go engine as of 2026, originally trained AlphaZero-style with extensive enhancements; supports score/territory prediction, handicap, multiple board sizes and rulesets via a single net. Leela Zero is the older AlphaGo-Zero replica, still widely deployed but typically weaker than KataGo. Both are command-line GTP engines requiring a GUI (KaTrain, Lizzie, LizzieYzy, q5Go, Sabaki). KataGo is famously vulnerable to adversarial attacks: Wang et al., "Adversarial Policies Beat Superhuman Go AIs" (arXiv:2211.00241, Nov 2022, ICML 2023): "We attack...KataGo by training adversarial policies against it, achieving a >97% win rate against KataGo running at superhuman settings." [GitHub + 2](#)

ZBaduk (zbaduk.com) — Browser-based, by Bram Vandenbon. Three tools: Smart Review (Leela Zero), KataGo full analysis, position analysis. Pricing approximately €4/month per Go Magic's tutorial. Stripe billing; ~80% of capacity reportedly serves free users. **Weaknesses:** smaller team than AI Sensei, less polished UI, periodic outages (SSL expired November 2024 per [webrate.org](#)), no NL commentary. [Zbaduk + 4](#)

BadukPop (CorePlane Inc.) — Mobile-first (iOS/Android). 5,000+ tsumego problems, lessons, AI opponents, online play, correspondence games. Subscription-gated practice mode. The vendor's own App Store reply: "Thanks so much for the review! We agree 100% with your feedback, and are working on more in-depth explanations for puzzle answers right now." App Store reviewer: "I wish it would explain why a move is wrong because I don't always understand the correct answer." — directly validating the explanation gap. [Uptodown + 2](#)

Crazy Stone DeepLearning Pro (UNBALANCE Corp.) — Standalone iOS/Android app, descended from the 2013-era Crazy Stone engine (Rémi Coulom). 17–20 levels of strength, analysis mode (move list, histogram, situation graph). Subscription required for ad-free + higher levels. Reaches ~5-dan KGS strength. Useful for offline mobile play; analysis is light compared to KataGo. [App Store](#) [App Store](#)

SmartGo (Anders Kierulf) — Long-standing commercial product (SmartGo for Windows ships with a 45,000+ pro-game database). Recently moved its iOS Kifu product to a subscription model, prompting some defection to Crazy Stone. SGF editing + database research + lightweight AI. Not primarily an AI reviewer. [AlternativeTo](#) [App Store](#)

Go Quest (Mindwalk Corp.) — Free, quick-game mobile app, 9×9/13×13/19×19, peak-hours real-time play, Tsumego Challenge mode. Casual market; no AI review feature. [Google Play](#) [Google Play](#)

Korean/Chinese platforms — Fox (Foxwq), Tygem, WBaduk (CyberOro), Pandanet (IGS): Fox is the most-trafficked Go server globally (Chinese audience), Tygem #2 (Korean audience), WBaduk Korean. Per a 2024 Medium overview by Go Magic: Tygem and WBaduk have "AI Analysis: N/A" in their feature matrices. Pandanet (IGS) explicitly integrates with AI Sensei. **Fox has integrated AI analysis** (used by professionals for review, and at the center of cheating scandals per CNN's 2023 reporting), but the Fox client is Chinese-language-only and not actively maintained for Western users. [Medium + 2](#)

DeepMind AlphaGo-derived tools — DeepMind retired AlphaGo after the 2017 Future of Go Summit and never released a consumer product. The legacy was an open data set (50 AlphaGo vs AlphaGo games) and the methodological inspiration for KataGo, ELF OpenGo (Facebook, now archived), Golaxy, FineArt (Tencent), HanDol, DolBaram. Tencent's PhoenixGo was open-sourced; Facebook's ELF OpenGo is archived. [Wikipedia](#) [Wikipedia](#)

Chess precedent — gold standard for comparison:

- **Chess.com's "Game Review with Coach Explanations"** launched with the September 1, 2022 membership restructure (per Chess.com support article: *"As of September 1, 2022, Premium plans have changed for new subscribers"*). Each move is classified Brilliant/Great/Best/Inaccuracy/Mistake/Blunder; the Virtual Coach provides a one-line summary plus (for Diamond) text explanations of *why* a move was best/missed. Diamond costs ~\$10/mo annualised (\$120/yr) or ~\$19.99/mo billed monthly. User complaints persist about explanation quality ("Hmm... this seems like a mistake to me" is a real failure mode), but the tier exists and is profitable. [Chess.com + 5](#)
- **Lichess** offers Stockfish analysis (free, no explanations); a community proposal "AI-Powered Game Analysis and Commentary Feature Proposal for Lichess" exists in the forum but has not been productized. The closest is **ChessCoach** ([chrisbutner/ChessCoach](#)), an academic-grade engine with neural NL commentary running as bot @PlayChessCoach on Lichess. Per the ChessCoach GitHub README: *"It plays chess with a rating of approximately 3450 Elo, which means it should usually beat even the strongest human players at 2850 Elo, and many other engines, but will often lose to the strongest, such as Stockfish 14 at 3550 Elo."* [GitHub](#)

Market Context

- **Total Go-playing population:** The often-cited 2002 Japanese census put the worldwide figure at ~24 million. The British Go Association's 2016 update suggested 60 million. The Chinese Weiqi Association says 40 million play in China alone, with 200,000 schools (per

CNN 2023). Per Google DeepMind's official AlphaGo page: *"AlphaGo's 4-1 victory in Seoul, South Korea, in March 2016 was watched by over 200 million people worldwide"* — a lower bound on Go-aware global audience. (Wikipedia + 2)

- **Geography:** China » Korea > Japan » Taiwan, then a long tail in Europe, North America, Vietnam, Russia, Ukraine. The British Go Association estimates "more than 70 countries where organised Go is played." (British Go Association)
- **Languages:** Chinese (Simplified), Korean, Japanese dominate; English is the lingua franca of Western Go and most of OGS; Russian, German, French communities also active.
- **Platforms by users:** Foxwq (largest, mostly Chinese, no good Western client), Tygem (~Korean, second largest), WBaduk (~6,000 average concurrent, Korean), Pandanet/IGS (~2,200 concurrent, Japanese), OGS (English/multilingual web), KGS (English-speaking, Java client, declining), Go Quest (mobile-casual). AI Sensei explicitly supports SGF imports from all of these. (Medium) (Medium)
- **Pain points** (consistently surfaced across forums):
 1. "The AI tells me my move was -8 points, but not why" — winrate/score loss is not a teaching signal at kyu levels.
 2. "Reading AI variations 10 moves deep is overwhelming when I don't understand the first move."
 3. "I review my game with AI, nod thoughtfully, then make the same mistake two games later" (paraphrasing AI Sensei's own blog).
 4. Setting up local AI (KataGo + GUI + GPU drivers) is friction for non-technical players.
 5. Mobile review is weak — only AI Sensei and ZBaduk are truly mobile-friendly via PWA; KaTrain/Lizzie are desktop-only.

Natural Language Go Commentary — The Gap

There is no shipping Go product that generates move-by-move natural-language explanations of Go play. This was confirmed via dedicated search of academic literature, GitHub, and commercial offerings:

- The Adarie article "LLMs Playing and Commentating on Go: Current State (2025)" states: *"A typical Go engine can output a lot of data about a position – e.g. 'win rate now 67% for Black', 'best move at E4, variation shows sequence ...' – but it does not speak in human-friendly terms. It won't tell you why a move is good in plain language or what strategy each player is following (at least not without additional tools)."* (Adarie)
- The MIT Technology Review article "AI is rewiring how the world's best Go players think" (Feb 27, 2026) quotes Korean pro Kim describing the AI's *"blinking screen showed the winning probabilities of each move, with no explanations. Even top players like Kim and Shin don't understand all of AI's moves."* (MIT Technology Review)
- Academic precedent in chess (Jhamtani et al. 2018 "Learning to Generate Move-by-Move Commentary"; Zang et al. 2019 "Automated Chess Commentator Powered by Neural Chess Engine") shows the technique is feasible — ChessCoach (Chris Butner, 2021) is a 3,450-Elo

engine that plays + comments in plain English on Lichess. Nothing analogous has been built for Go.

- The 2020 arXiv paper "The Go Transformer: Natural Language Modeling for Game Play" (Ciolino, Kalin, Noever) used GPT-2 to generate SGF *moves*, not commentary — so the closest Go-LLM crossover work is in move generation, not explanation.
- All OGS forum threads asking for "what do the numbers mean" / "why is this move bad" remain unaddressed at the product level: e.g., user Jocaramiro, July 9, 2019: *"It would be helpful to get a pointer to where Leela explains, or has explained for them, what it is doing. I saw elsewhere that the preferred move was highlighted in blue. What do the numbers mean?"* (Online Go Forum)

Implication for product strategy: An LLM layer on top of KataGo (taking principal variations + score deltas + ownership maps and translating them into prose like "Black should have approached the 4-4 stone from the open side because the lower-right corner is already settled, and pressing here builds a moyo facing the right side") would be a category-defining feature. The technical building blocks (KataGo Python API, OpenAI/Anthropic/Gemini APIs, board → text prompt engineering) are all commodity. The closest analog in another mind-sport is Chess.com's Diamond Coach.

Community Sentiment — Direct Quotes

- **OGS forum, n0w3l, Feb 18, 2021:** *"So I would like to kindly ask for anyone to review any game with 'full AI analysis' with the mentality of teaching newcomers how to review games using OGS's AI further than 'move X made me loose 80% winrate' ... PS1- I'm aware of the limitations of non-explainable AI's, I'm just looking for pointers about the available tools."* (online-go)
- **OGS forum, "Request: AI review of variations":** *"I'm reviewing an AI variation, and don't understand why move 6 is what looks like an unnecessary defensive move. I would love to be able to play the move I think is correct, in a variation, and have the AI evaluate that..."* (Online Go Forum)
- **AI Sensei blog (self-aware about the gap):** *"You play a game of Go. You review it with AI. You see where you went wrong. You nod thoughtfully, close the review, and then... make the exact same mistake two games later."* (Ai-sensei)
- **Life in 19x19, hypersleep, "How to review your own games as a kyu?":** *"I usually look at my games with AI and look for the moves that turned the tables. I don't understand most of the moves the AI recommends, so I don't delve into the variations."* This is consistent with senior responders telling kyu-level players to ignore engine variations they can't parse — a tacit admission of the explanation gap. (Life In 19x19)
- **App Store, BadukPop user:** *"I wish it would explain why a move is wrong because I don't always understand the correct answer."* The developer publicly replied: *"We agree 100% with your feedback, and are working on more in-depth explanations for puzzle answers right now."* (App Store)
- **MIT Tech Review, pro Kim:** *"It seems like it's thinking in a higher dimension."* This frames the cultural problem — AI is opaque even to professionals. (MIT Technology Review)

- **OGS forum (re: AI Sensei comparison):** "OGS offers reviews of the games, but they do not allow you to create records of where you made mistakes so that you can quiz yourself and review the mistakes and then make the right move. The paid level of both is about \$5.00 USD per month. But you don't just get told where your moves were poor: you are provided a means to turn those poor moves into flashcards that get reviewed on a regular basis." (Players are willing to pay ~\$5/mo for review tooling.) [Online Go Forum](#)

OGS API and Extension Ecosystem

- **APIs available:** A REST API documented at apidocs.online-go.com and apiary.io/ogs.docs (covers logging in, fetching games, SGFs, user data) and a Socket.IO real-time channel for live game state, chat, and move events. Useful tutorial: "OGS API notes" forum thread by community member.
- **Existing third-party userscripts on Greasy Fork** (all cosmetic, all low-install):
 1. OGS Utilities (Tomasz Błajek, 2015, 24 installs) — move counter in browser title. [greasyfork](#)
 2. OGS max vertical board (Brendan R, 2017, 19 installs). [greasyfork](#)
 3. Remove OGS Nav Logo (metaperl, 2023, 14 installs). [greasyfork](#)
 4. Copy SGF from OGS (Pavel Watson, 2024, 3 installs). [greasyfork](#)
 5. scroll-navigate-ogs (TransientError, 2024, 1 install). [greasyfork](#)
 6. Goban Resizer (Hesper Croft, 2025, 3 installs). [greasyfork](#)
 7. OGS Game History Widener (BurrowBird, 2026, 1 install). [greasyfork](#)
- **Browser extensions:** OGS Notifier ([prozz/ogs-notifier](#), forked by traveller42 on Chrome Web Store) — sends desktop notifications for waiting games. Periodically broken by OGS updates. [GitHub](#) [Online Go Forum](#)
- **No third-party commentary, review, or LLM-integration extension has ever shipped for OGS.** The official online-go GitHub organization has ~10 repos, none of them third-party reviewer overlays. [GitHub](#)
- **Implication:** A Chrome/Firefox extension that overlays AI-Sensei-like or LLM-narrated commentary onto OGS live and finished-game pages is a clear, technically tractable wedge. The Chrome 138+ per-extension "Allow User Scripts" toggle simplifies userscript distribution at scale. [Chrome Developers](#)

Recommendations

Stage 1 (build now): Ship a minimum-viable LLM-narrated Go review wrapping KataGo. Take the KataGo analysis JSON (winrate, score-lead, ownership map, top-5 moves, principal variations) and prompt a frontier LLM (GPT-4-class or Claude) to produce 1-3 sentences per key move explaining the strategic reason in human Go vocabulary (influence, territory, aji, miai, sente, shape, direction of play). This is the category-defining feature no competitor ships. Target English first; Korean, Chinese, Japanese later.

Stage 2 (distribute): Build it as a browser extension that overlays onto OGS finished-game pages — the install funnel is wide-open (no competing extension; tens of thousands of OGS users), the API is documented, and you piggyback on OGS's free Leela analysis as a fallback before paying for KataGo cycles. Adopt a freemium model anchored at AI Sensei's known ~\$5/mo Basic-tier price point (Chess.com's Diamond at ~\$10/mo annualised is the upper anchor); commodity LLM cost per game is well under \$0.10 with prompt caching.

Stage 3 (mobile + integrate): Once English NL commentary has product-market fit, follow AI Sensei's playbook and partner with Pandanet (Japan), then negotiate Korean (Tygem) and Chinese (Foxwq) integrations. Mobile PWA + native iOS/Android app come next; the BadukPop user base and SmartGo refugees are addressable.

Benchmarks that should shift the strategy:

- If >5% of OGS finished-game viewers install the extension within 6 months, double down on a standalone web product (compete head-on with AI Sensei).
- If retention is poor on free tier, pivot to spaced-repetition flashcard generation (the user-quoted advantage that drove some users to AI Sensei: *"you are provided a means to turn those poor moves into flashcards that get reviewed on a regular basis"*). ([Online Go Forum](#))
- If OGS itself ships LLM commentary (anoek's team has been responsive to community asks), pivot to the underserved Korean/Chinese platforms where built-in AI review is "N/A."
- If LLM unit economics break above \$0.50/review, switch to local-model deployment (Llama 3 70B class) on a single GPU server pooled across users.

Threats to watch:

- The Chess.com / Lichess pattern suggests the official OGS team could ship Coach Explanations themselves. Anoek's team is small but technically able.
- AI Sensei could add LLM commentary as a Pro/Dan tier feature. Their KataGo backend is the same, so they'd ship faster.
- Frontier LLMs may add "Go expert mode" as part of broader training (unlikely soon, but possible).

Caveats

- **Player-count figures vary wildly** (24M – 60M+ players; 200M+ AlphaGo viewers) and are based on dated censuses or single-event viewership. The actively-paying English-speaking Go review tool market is much smaller — likely tens of thousands to low hundreds of thousands. No public, recent, audited number exists.
- **AI Sensei's subscriber count is not disclosed**, so its revenue baseline is inferred from pricing and community references. Founder team is small (4 named developers).
- **Fox / Tygem / WBaduk AI features:** Some sources say Fox has AI analysis integrated; others say "N/A." Chinese-language documentation is the authority and was not deeply searched here. The Fox client is hostile to non-Chinese users.

- **Chess.com "Coach Explanations" quality is uneven** — multiple user complaints document failure modes like "Hmm... this seems like a mistake to me" with no further reasoning. An LLM-on-KataGo product for Go could face the same hallucination/quality risk; thorough evaluation against pro reviewers is essential.
- **Cheating risk:** Any browser-extension AI overlay is fundamentally a cheating tool if it runs during live games. The CNN-reported AI-cheating scandals in Korean pro Go (2022–2023) make this an existential reputational risk; the product **MUST** restrict commentary to finished games / review mode only, or run with explicit opponent consent for live teaching games. (CNN)
- **Engine "blind spots":** KataGo is famously vulnerable to adversarial cyclic-group attacks — Wang et al., "*Adversarial Policies Beat Superhuman Go AIs*" (arXiv:2211.00241; ICML 2023) achieved ">97% win rate against KataGo running at superhuman settings." Commentary built on KataGo inherits these blind spots and should never be presented as infallible — clear UX framing is required. (arXiv)
- **Forum quotes are dated** (some from 2019–2021); a deeper sweep of 2024–2026 Reddit r/baduk threads was not completed because of search-budget limits and may surface newer dissatisfaction signals.